

M Sai Vipul Chinmay

Hyderabad, India | +91 83743 83285 | m.sai.vipul.18@gmail.com | linkedin.com/in/vipulchinmay

PROFESSIONAL SUMMARY

A Computer Science undergraduate with hands-on experience designing and deploying ML/NLP systems. AI Engineer Intern at Krignal Technologies, building LLM pipelines, multi-agent AI architectures, RAG systems, and ranking-aware retrieval services. Strong foundation in Python, Java, and SQL. Deep interest in search relevance, learning-to-rank, and large-scale ML model deployment.

TECHNICAL SKILLS

Languages: Java, Python, C++, JavaScript, SQL, C

ML / AI: LLM Fine-Tuning, RAG Architecture, Multi-Agent AI, LangChain, LangGraph, Transformers (HuggingFace), scikit-learn, LightGBM, NLP, Model Evaluation (NDCG, MRR, Precision@K), Langchain, LangGraph, Information Retrieval, BM25, TF-IDF, Learning-to-Rank (LTR), Vector Search (FAISS), Semantic Search, Query Understanding, A/B Testing.

Cloud & Infra: AWS, Azure, Docker, kubernetes, CI/CD, Jenkins, MySQL, MongoDB, DynamoDB, Redis, Grafana.

Web / Backend: Node.js, SpringBoot, React.js, REST API Design, MERN Stack

PROFESSIONAL EXPERIENCE

AI Engineer Intern | Krignal Technologies

Nov 2025 – Present | Hyderabad, India

- Collaborating with AI and Product teams to design, build, and test AI-driven backend services powering content generation, tagging, access control, and real-time insight delivery across professional communities.
- Working on a multi-agent AI architecture where specialized agents handle summarization, forecasting, opportunity detection, and contextual insights; contributing to frontend dashboards, personalized insight feeds, expert-member interfaces, and match request workflows.
- Assisting in user and access management systems, implementing role-based permissions and secure onboarding flows for experts, members, and administrators; participating in QA, testing, and performance optimization for scalable AI outputs across devices.
- Supporting AI agent evaluation and prompt experimentation, optimizing model behavior across different personas, tone styles, and community contexts using RAG pipelines and fine-tuned transformer models.

PROJECTS

Search Ranking Pipeline: BM25 + Neural Reranker | *LightGBM, FAISS, sentence-transformers, FastAPI, AWS* 2025

- Implemented a two-stage retrieval system — BM25 for candidate generation, cross-encoder reranker for relevance scoring — evaluated on the Amazon ESCI dataset.
- Trained a LightGBM learning-to-rank model incorporating semantic features from sentence-transformers; achieved 12% NDCG@10 improvement over BM25 baseline.
- Designed offline A/B evaluation framework to compare ranking approaches; deployed serving endpoint via FastAPI on AWS with sub-200ms P95 latency.

Seekhan — RAG-powered AI Quiz Generator | *LangChain, FAISS, HuggingFace, React.js, Node.js*

2024–2025

- Designed document ingestion and chunking pipeline with FAISS-indexed sentence-transformer embeddings for semantic retrieval.
- Built structured output generation using LLMs for MCQ, True/False, and Fill-in-the-blank formats with factual consistency evaluation.
- Engineered natural language query interface for educators to generate tailored quizzes through conversational prompts.

AushadX — Multilingual Medicine Assistance App | *MongoDB, Express.js, React.js, Node.js, Easy-OCR, Docker* 2024

- Built full-stack MERN application with Easy-OCR pipeline for automatic extraction of medicine names, dosage, and expiry data from images.
- Integrated multilingual NLP pipeline (5+ languages) and a fine-tuned medical QA model; deployed on AWS with Docker containerization.
- Integrated geolocation API to surface nearby hospitals; built health analytics dashboard for medication tracking.

ADDITIONAL INFORMATION

Languages: English (Fluent), Hindi (Fluent), Telugu (Native)

Interests: Music, Cricket, Movies, LLM systems, blockchain technology, emerging AI trends.